

Computer Vision HW2 Report

Student ID: R10943019


Name: 李奕勳

Part 1. (10%)


- Plot confusion matrix of two settings. (i.e. Bag of sift and tiny image representation) (5%)

Ans:

Tiny image



Bag of sift



- Compare the results/accuracy of both settings and explain the result. (5%)

Ans:

Tiny image:

```
$ python3 p1.py --feature tiny_image --classifier nearest_neighbor
Getting paths and labels for all train and test data
knn calculating.....
Accuracy = 0.234
```

只透過 resize 圖片來當成特徵，無法有效地得到有意義的特徵表示，即使是同一類別的圖片，彼此之間還是有很多變異，僅透過 tiny image 當作特徵來分群，很難將同一類別的圖片聚集在一起，因此在 testing 時，準確率自然就不會高，但優點是每張圖片只有一個特徵向量，因此運算速度很快。

Bag of sift:

```
$ python3 p1.py --feature bag_of_sift --classifier nearest_neighbor
Getting paths and labels for all train and test data
(1, 400)
knn calculating.....
Accuracy = 0.604
```



SIFT 能找到圖片中不同維度的特徵，例如:edge、corner 等等.....，透過 SIFT 得到的特徵點相較於 tiny image 更有意義，因此就能更容易地將同類別的圖片分群在一起，在 testing 時，準確率就相對的高上許多，但缺點是同一張圖片中，SIFT 會找到很多特徵點，若將所有特徵點納入考慮，運算複雜度會很高。

Part 2. (35%)

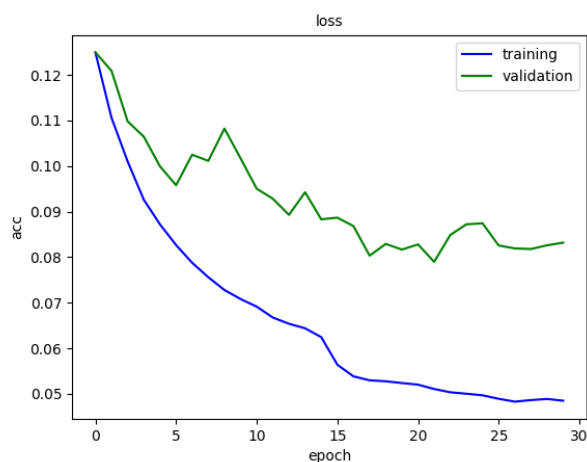

• Compare the performance on residual networks and LeNet. Plot the learning curve (loss and accuracy) on both training and validation sets for both 2 schemes. 8 plots in total. (20%)

Ans:

LeNet



residual networks



從圖中可以發現，無論是 training 或是 validation，在 accuracy 的表現上都是 residual network 較好，透過 residual block 可以避免梯度消失的問題，因此相較於 LeNet，residual network 較可以建立起深層的網路來學習不同語意的特徵，在分類問題上得到更好的準確率。

- Attach basic information of the model you use including model architecture and number of the parameters. (5%)

Ans:

Model I use : EfficientNetB0 pretrained on ImageNet

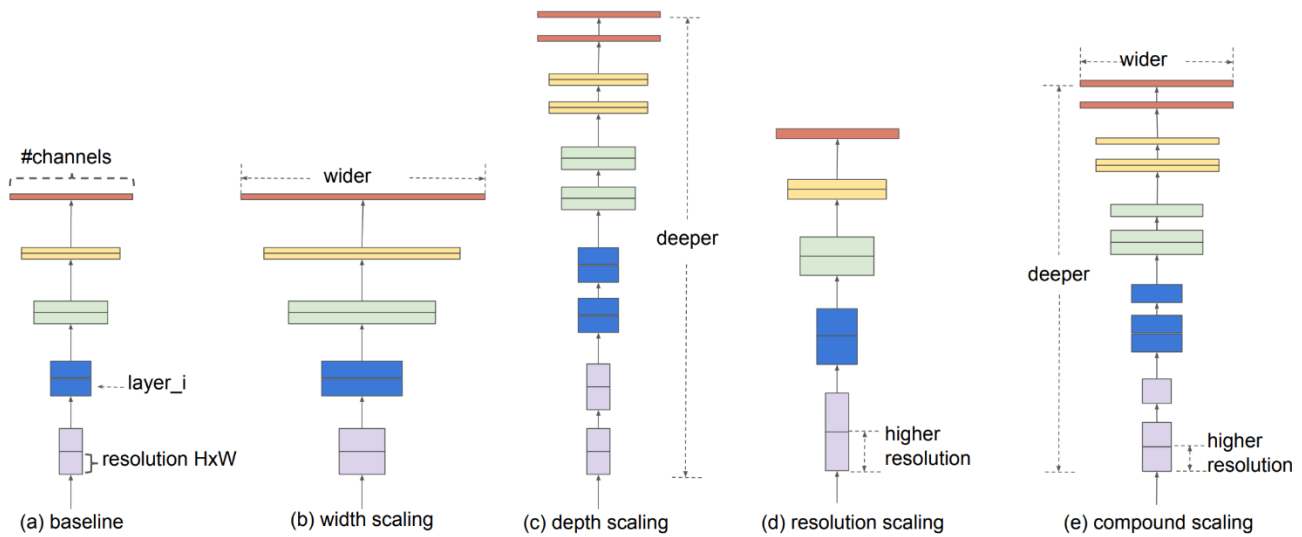


Figure 2. **Model Scaling.** (a) is a baseline network example; (b)-(d) are conventional scaling that only increases one dimension of network width, depth, or resolution. (e) is our proposed compound scaling method that uniformly scales all three dimensions with a fixed ratio.

Ref. [EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks \(arxiv.org\)](https://arxiv.org/abs/1907.11164)

Key idea : compound scaling

在神經網路的 scaling 有三個維度:深度、寬度、解析度

在 EfficientNet 中提到 compound scaling，必須平衡三個維度的 scaling 幅度，才能得到更有效率且準確率高的網路。

Model architecture for EfficientNetB0:

└─MBConvBlock: 2-14	--
└─Conv2dStaticSamePadding: 3-107	221,184
└─BatchNorm2d: 3-108	2,304
└─Conv2dStaticSamePadding: 3-109	28,800
└─BatchNorm2d: 3-110	2,304
└─Conv2dStaticSamePadding: 3-111	55,344
└─Conv2dStaticSamePadding: 3-112	56,448
└─Conv2dStaticSamePadding: 3-113	221,184
└─BatchNorm2d: 3-114	384
└─MemoryEfficientSwish: 3-115	--
└─MBConvBlock: 2-15	--
└─Conv2dStaticSamePadding: 3-116	221,184
└─BatchNorm2d: 3-117	2,304
└─Conv2dStaticSamePadding: 3-118	28,800
└─BatchNorm2d: 3-119	2,304
└─Conv2dStaticSamePadding: 3-120	55,344
└─Conv2dStaticSamePadding: 3-121	56,448
└─Conv2dStaticSamePadding: 3-122	221,184
└─BatchNorm2d: 3-123	384
└─MemoryEfficientSwish: 3-124	--
└─MBConvBlock: 2-16	--
└─Conv2dStaticSamePadding: 3-125	221,184
└─BatchNorm2d: 3-126	2,304
└─Conv2dStaticSamePadding: 3-127	28,800
└─BatchNorm2d: 3-128	2,304
└─Conv2dStaticSamePadding: 3-129	55,344
└─Conv2dStaticSamePadding: 3-130	56,448
└─Conv2dStaticSamePadding: 3-131	221,184
└─BatchNorm2d: 3-132	384
└─MemoryEfficientSwish: 3-133	--
└─MBConvBlock: 2-17	--
└─Conv2dStaticSamePadding: 3-134	221,184
└─BatchNorm2d: 3-135	2,304
└─Conv2dStaticSamePadding: 3-136	10,368
└─BatchNorm2d: 3-137	2,304
└─Conv2dStaticSamePadding: 3-138	55,344
└─Conv2dStaticSamePadding: 3-139	56,448
└─Conv2dStaticSamePadding: 3-140	368,640
└─BatchNorm2d: 3-141	640
└─MemoryEfficientSwish: 3-142	--
└─Conv2dStaticSamePadding: 1-4	--
└─Identity: 2-18	--
└─BatchNorm2d: 1-5	2,560
└─AdaptiveAvgPool2d: 1-6	--
└─Dropout: 1-7	--
└─Linear: 1-8	12,810
└─MemoryEfficientSwish: 1-9	--

To release GPU memory

Number of parameters for EfficientNetB0:

```
=====  
Total params: 3,609,894  
Trainable params: 3,609,894  
Non-trainable params: 0  
=====
```

Model architecture for EfficientNetB1:

Layer (type:depth-idx)	Param #
Conv2dStaticSamePadding: 1-1	--
└─ZeroPad2d: 2-1	--
BatchNorm2d: 1-2	64
ModuleList: 1-3	--
└─MBConvBlock: 2-2	--
└─Conv2dStaticSamePadding: 3-1	288
└─BatchNorm2d: 3-2	64
└─Conv2dStaticSamePadding: 3-3	264
└─Conv2dStaticSamePadding: 3-4	288
└─Conv2dStaticSamePadding: 3-5	512
└─BatchNorm2d: 3-6	32
└─MemoryEfficientSwish: 3-7	--

To release GPU memory

- .
- .

23 MBConvBlock in a series

- .
- .

└─MBConvBlock: 2-24	--
└─Conv2dStaticSamePadding: 3-195	614,400
└─BatchNorm2d: 3-196	3,840
└─Conv2dStaticSamePadding: 3-197	17,280
└─BatchNorm2d: 3-198	3,840
└─Conv2dStaticSamePadding: 3-199	153,680
└─Conv2dStaticSamePadding: 3-200	155,520
└─Conv2dStaticSamePadding: 3-201	614,400
└─BatchNorm2d: 3-202	640
└─MemoryEfficientSwish: 3-203	--
Conv2dStaticSamePadding: 1-4	--
└─Identity: 2-25	--
BatchNorm2d: 1-5	2,560
AdaptiveAvgPool2d: 1-6	--
Dropout: 1-7	--
Linear: 1-8	12,810
MemoryEfficientSwish: 1-9	--

Number of parameters for EfficientNetB1:

Total params: 6,115,530
Trainable params: 6,115,530
Non-trainable params: 0

• Briefly describe what method do you apply? (e.g. data augmentation, model architecture, loss function, semi-supervised etc.) (10%)

Ans:

Data augmentation :

TOO MUCH augmentation will lead to poor training progress!!

```
transforms.RandomPerspective(distortion_scale = 0.6, p = 1.0),
transforms.RandomRotation(degrees = (0, 180)),
transforms.RandomAffine(degrees = (30, 70), translate = (0.1, 0.3), scale = (0.5, 0.75)),
transforms.RandomHorizontalFlip(p=0.5),
transforms.RandomVerticalFlip(p=0.5),
```

```
epoch = 13
100%|████████████████████████████████████████████████████████████████████████████████| 1294/1294 [01:35<00:00, 13.52it/s]
100%|████████████████████████████████████████████████████████████████████████████████| 144/144 [00:01<00:00, 77.04it/s]
*****
time = 1.0000 MIN 37.6128 SEC, total time = 22.0000 Min 57.8580 SEC
training loss : 0.1230  train acc = 0.2817
val loss : 0.1307  val acc = 0.2300
=====

epoch = 14
100%|████████████████████████████████████████████████████████████████████████████████| 1294/1294 [01:30<00:00, 14.36it/s]
100%|████████████████████████████████████████████████████████████████████████████████| 144/144 [00:01<00:00, 76.01it/s]
*****
time = 1.0000 MIN 32.0195 SEC, total time = 24.0000 Min 29.8777 SEC
training loss : 0.1222  train acc = 0.2911
val loss : 0.1273  val acc = 0.2587
=====
```

Final data augmentation - only vertical and horizontal flip

```
transforms.RandomHorizontalFlip(p=0.5),
transforms.RandomVerticalFlip(p=0.5),
```



```
epoch = 13
100%|████████████████████████████████████████████████████████████████████████████████| 1294/1294 [01:42<00:00, 12.56it/s]
100%|████████████████████████████████████████████████████████████████████████████████| 144/144 [00:02<00:00, 62.00it/s]
*****
time = 1.0000 MIN 45.3177 SEC, total time = 25.0000 Min 27.8183 SEC
training loss : 0.0588  train acc = 0.6832
val loss : 0.0608  val acc = 0.6639
=====

epoch = 14
100%|████████████████████████████████████████████████████████████████████████████████| 1294/1294 [01:43<00:00, 12.55it/s]
100%|████████████████████████████████████████████████████████████████████████████████| 144/144 [00:02<00:00, 62.41it/s]
*****
time = 1.0000 MIN 45.3926 SEC, total time = 27.0000 Min 13.4020 SEC
training loss : 0.0575  train acc = 0.6842
val loss : 0.0591  val acc = 0.6700
=====
```



Model architecture :

Except from LeNet and residual block, I also train on EfficientNet family (B0 、B1 、B4)


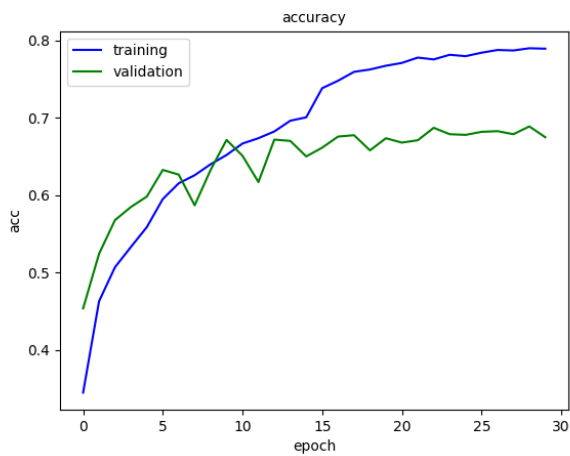
EfficientNetB0:



EfficientNetB1:



EfficientNetB4:



由上圖發現，如果使用同一個 training dataset 去訓練，越深層(B4)的網路越容易 overfitting

Loss function : use cross-entropy for classification problem